flytbase · bdr inbound assignment · walkthrough

Inbound Lead Desk

An AI agent that reads inbound leads from the inbox and produces a BDR’s first hour of work: qualification, cited research, a reply sequence, a case study, a partner motion and an AE handoff.

Samruddhi Sontakke claude code agent · gmail mcp · python mcp · mongodb · web app

samruddhisontakke.com · sontakkesamruddhijain89@gmail.com · +91 72196 00321 · Pune, India

the brief

One raw inbound email in. Six BDR artifacts out, written by a system.

01
Qualification
Framework, knowns, gaps, priority score

02
Deep account research
Real public data, cited, nothing invented

03
Response generation
Adaptive sequence with progression logic

04
Case study matching
From flytbase.com, referrer not hard-coded

05
Partner identification
Direct AE or partner-led, justified

06
AE handoff
Actionable next step with an owner

“You are not writing the answers. You are building the thing that writes the answers.” Every output in this video was produced by the system, and every requirement above is checked in code.

architecture

How the pieces connect

intake

Gmail

Gmail MCP: read and label only. Send is blocked.

→

the agent

Claude Code

Cron every 5 min runs /process-inbox and orchestrates

→

7 subagents

One per stage

Each a versioned file with its own prompt and schema

→

leaddesk mcp

Python tools

Scoring, sources, crawls, fact-check removal, PDF checks

MongoDB

Leads, every source with its raw text, inbox verdicts, agent heartbeat

Lead Desk web app

Live progress, a tab per stage, sources, agent trace. Never calls AI, never sends mail

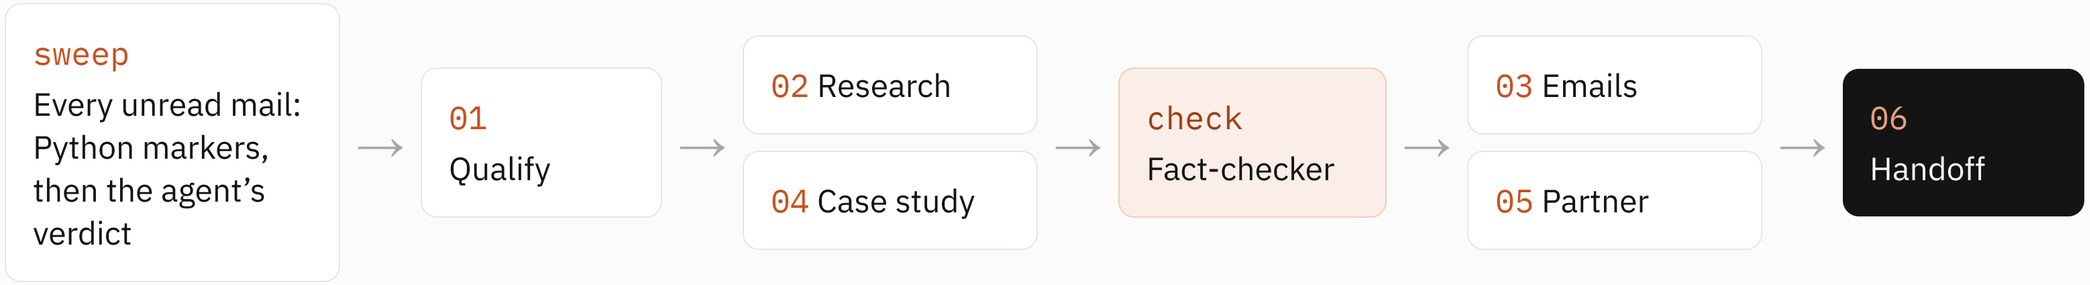
Research tools

WebSearch, a page and PDF reader, Reddit archive, flytbase.com sitemap

The rule: the AI reads, judges and writes. Python decides anything that must be exact.

one agent tick

Dependency-aware, not one long chain



Parallel where possible

Research and case study both need only Stage 1, so they run at once. Partner starts as soon as the case study is in.

Saved through Python

Each stage calls save_stage: validated, scored, checked against the PDF, stored in MongoDB and a run folder.

One lead at a time

A lease lock in MongoDB, Gmail message IDs deduplicated, and every subagent call written to a trace.

live 1 · switch to the system

The lead arrives in Gmail

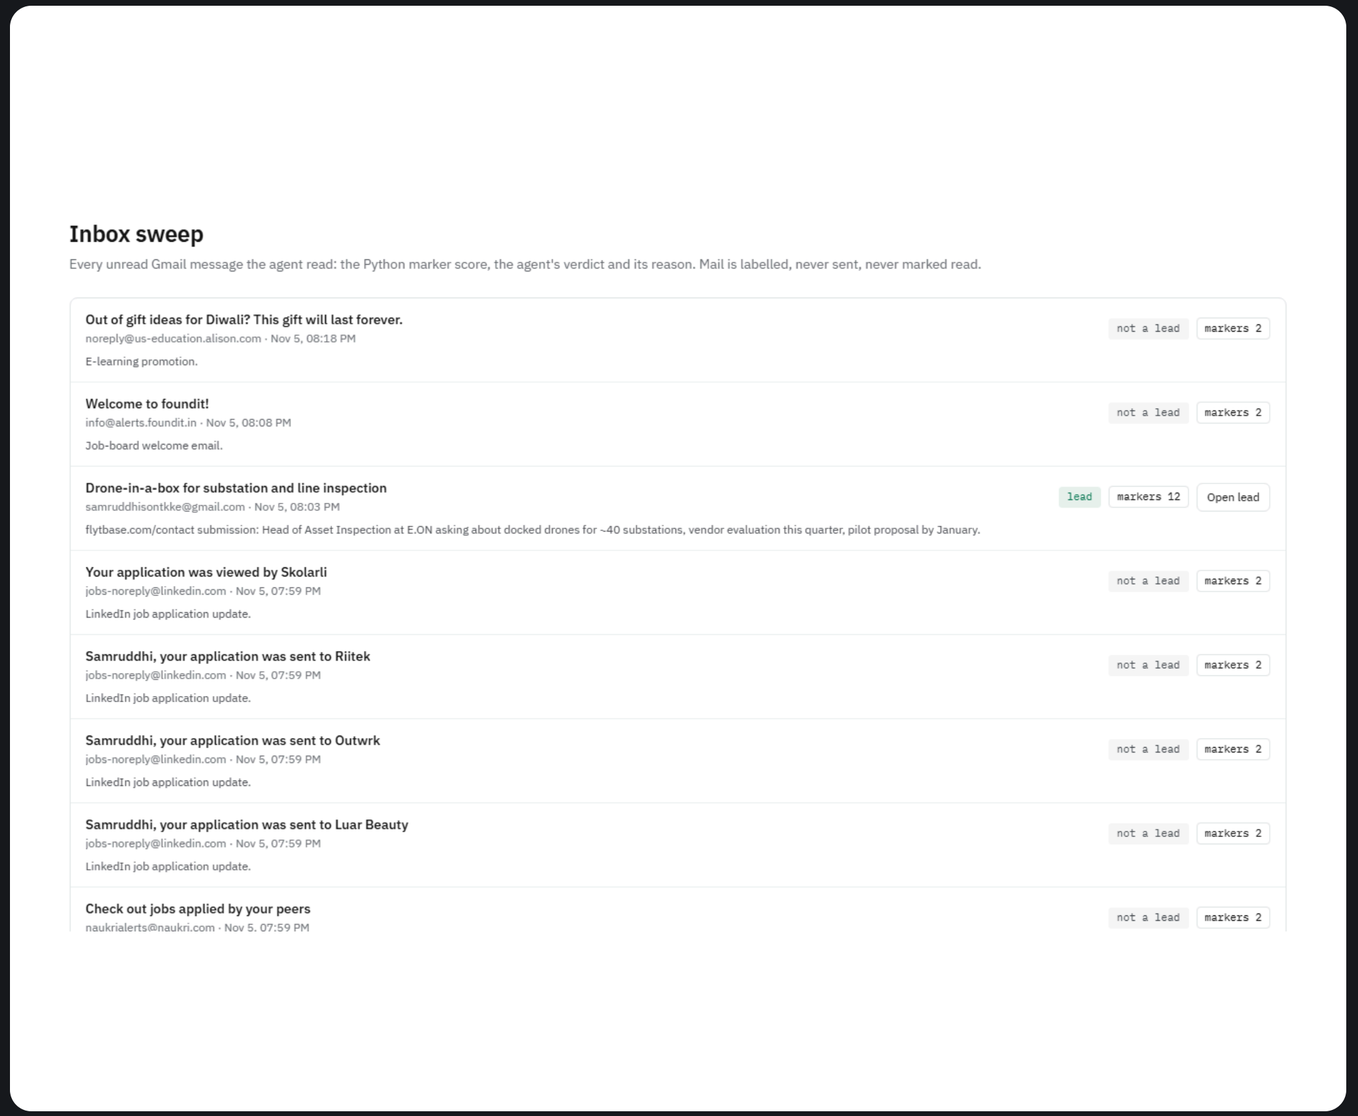
The agent swept 54 unread mails, mostly job alerts. Python scored each one and the agent gave a verdict with a reason. Only the SQM contact-form email was a lead (marker score 12).

gmail, via the gmail mcp

Autonomous inspection for our Atacama lithium sites
LeadDesk/Seen · LeadDesk/Lead · still unread

Drone-in-a-box for substation and line inspection
LeadDesk/Seen · LeadDesk/Lead · E.ON test

show: Gmail · /process-inbox · Inbox page · lead progress



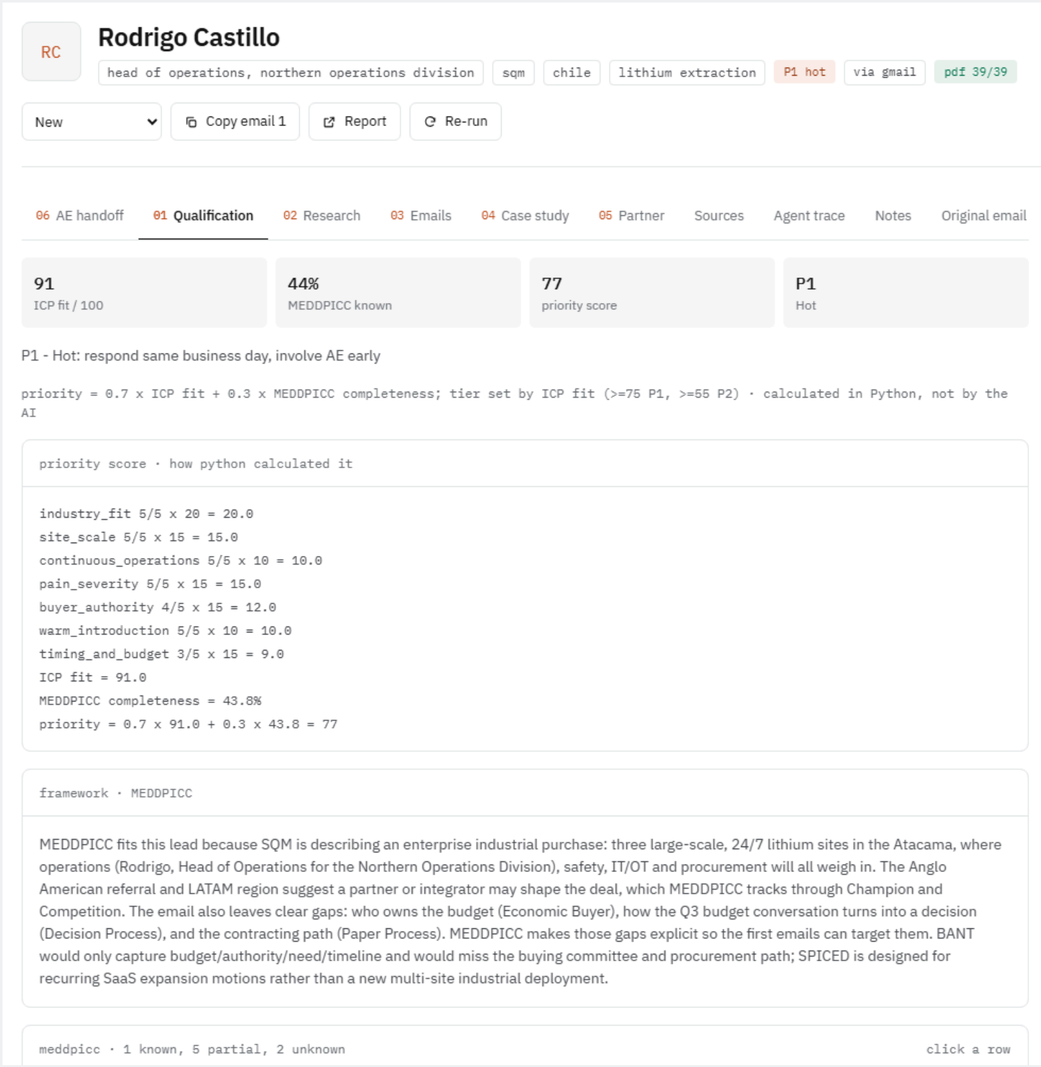
stage 01 · qualification · framework choice

MEDDPICC, scored in Python

91	44%	77
ICP fit	MEDDPICC	P1 hot

Why MEDDPICC: a multi-site industrial deal with a buying committee, procurement and a partner. BANT is too shallow, and SPICED suits SaaS expansion.

The AI only rates 7 fit factors with evidence. Python computes the score and shows its working.



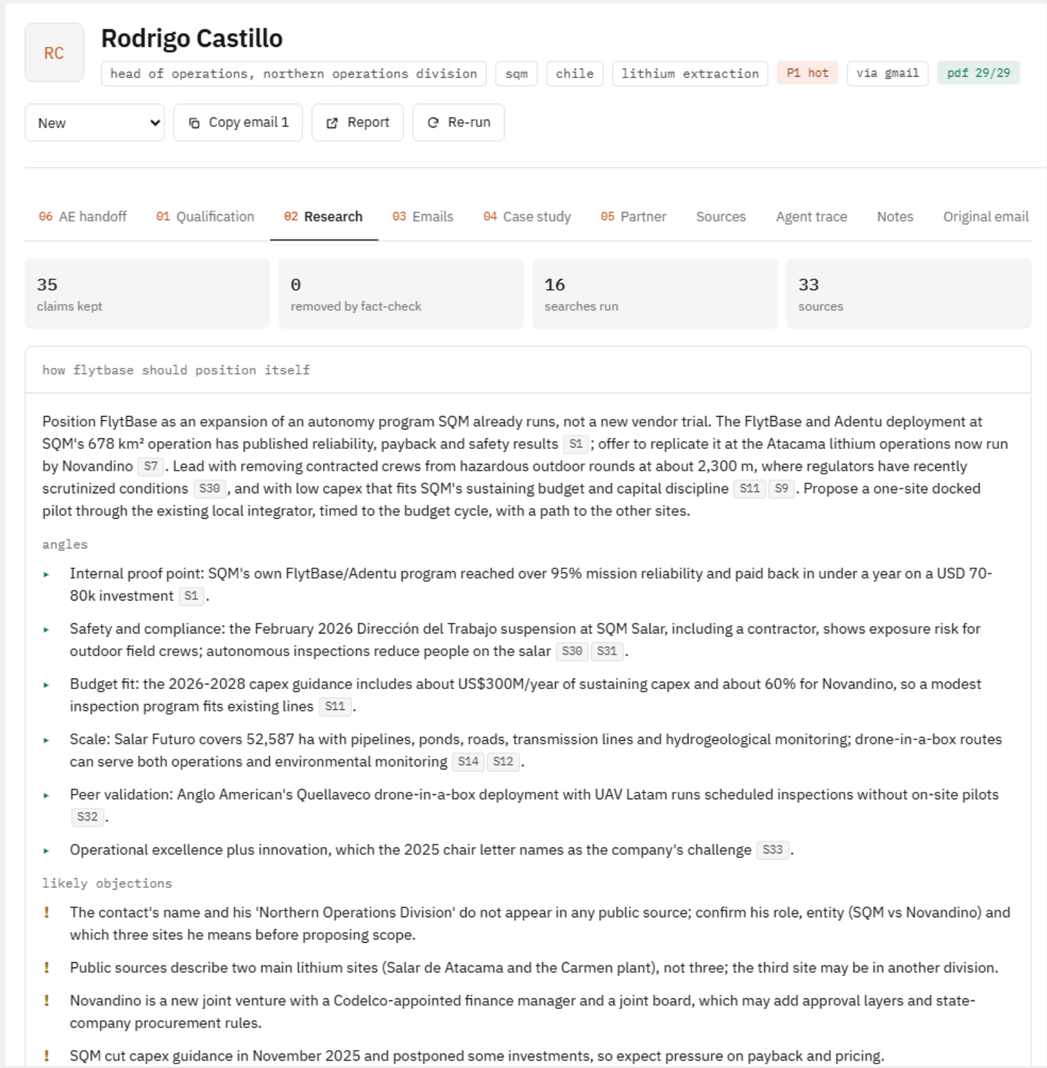
stage 02 · deep account research

Real sources, then an independent fact-check

33	41/43
sources	supported

Org structure, budget signals, news, investor letters, positioning. Every page is stored with its raw text as S#, and claims may only cite those IDs.

Key finding: SQM’s Atacama lithium now runs under Novandino, a Codelco–SQM joint venture, so the buyer may not be SQM S.A.



stages 04 + 05 · case study and partner

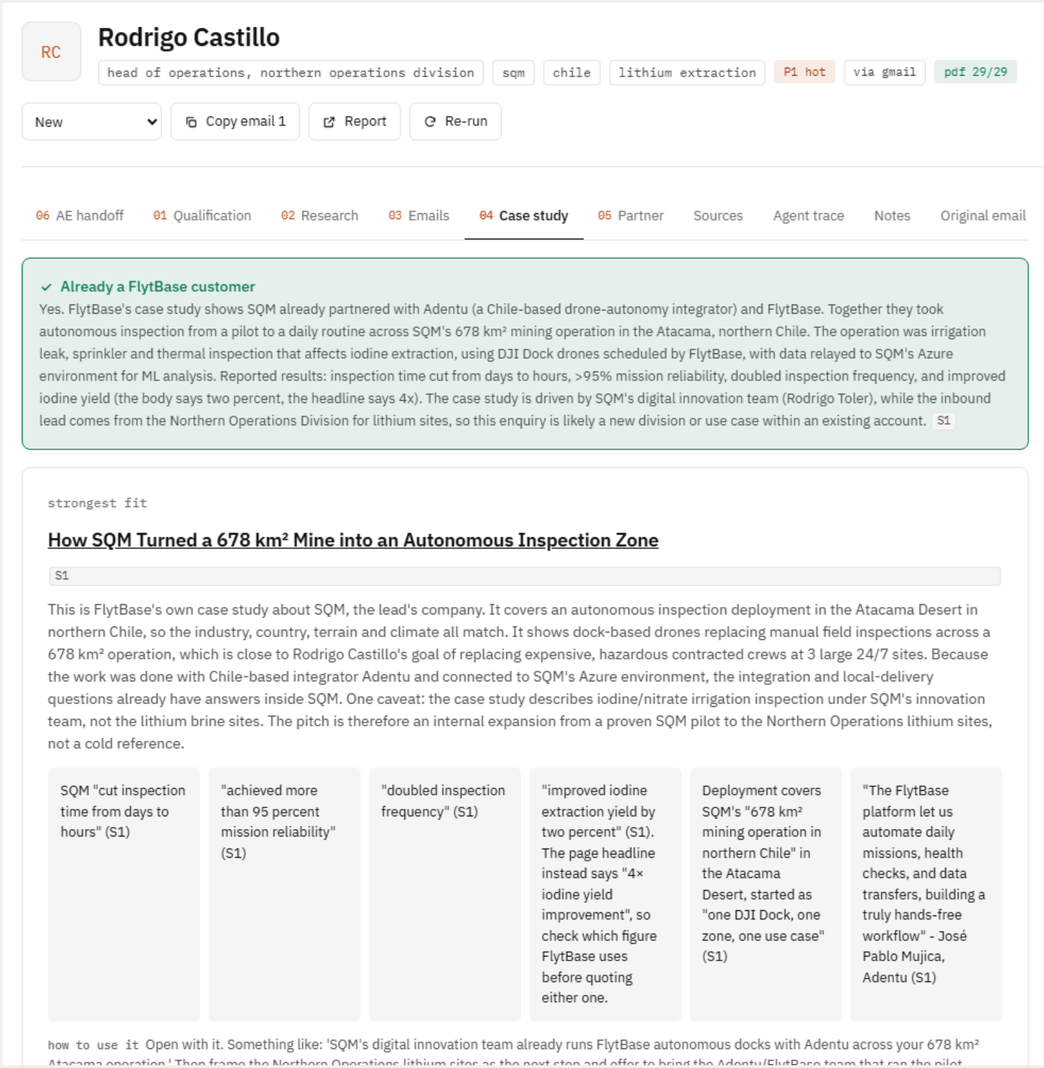
SQM is already a FlytBase customer

The crawl read all 24 case studies on flytbase.com. Python flagged the ones naming the lead's company or the referrer. Nothing is hard-coded.

Primary: SQM's own 678 km² deployment with Adentu in the Atacama

Referral link: Anglo American, Quellaveco in Peru, with UAV Latam

Motion: AE-led co-sell with Adentu; Dronestore Chile as backup



stage 03 · response generation

A reply to an inbound, not a cold email

Written from a playbook I built from Seth Godin, Alex Berman, Laura Lopuch and the marketingskills library, with every rule sourced.

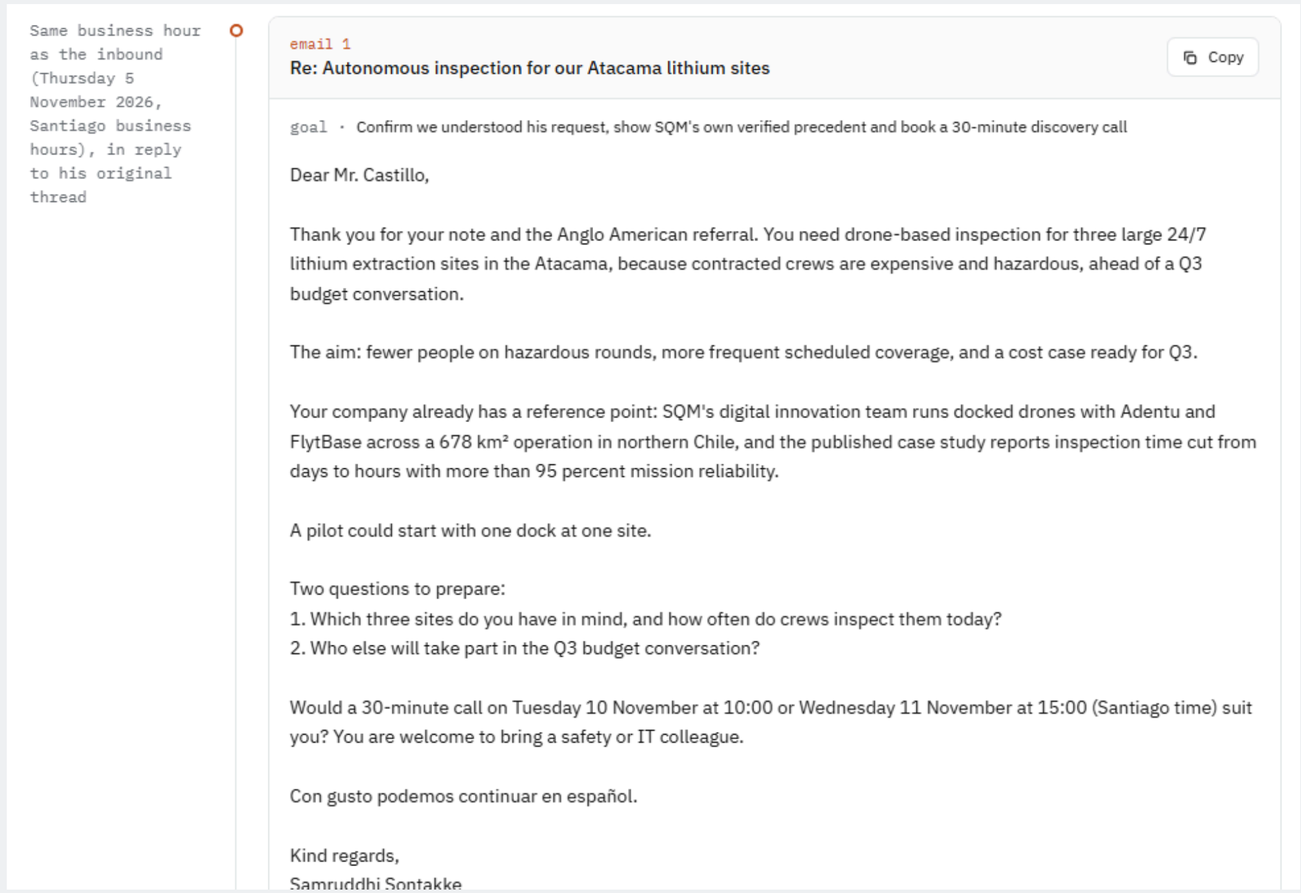
1 · same hour: mirrors his request, one proof point, 2 questions, 2 Santiago-time slots, Spanish offer

2 · Nov 10: IT and connectivity precedent → decision criteria

3 · Nov 18: Anglo American peer case + a forwardable summary

4 · Nov 25: closes the loop, offers to hold until Q3

Python lints 10 playbook rules. The first draft of email 1 was 212 words, so it was rejected and trimmed to 192.



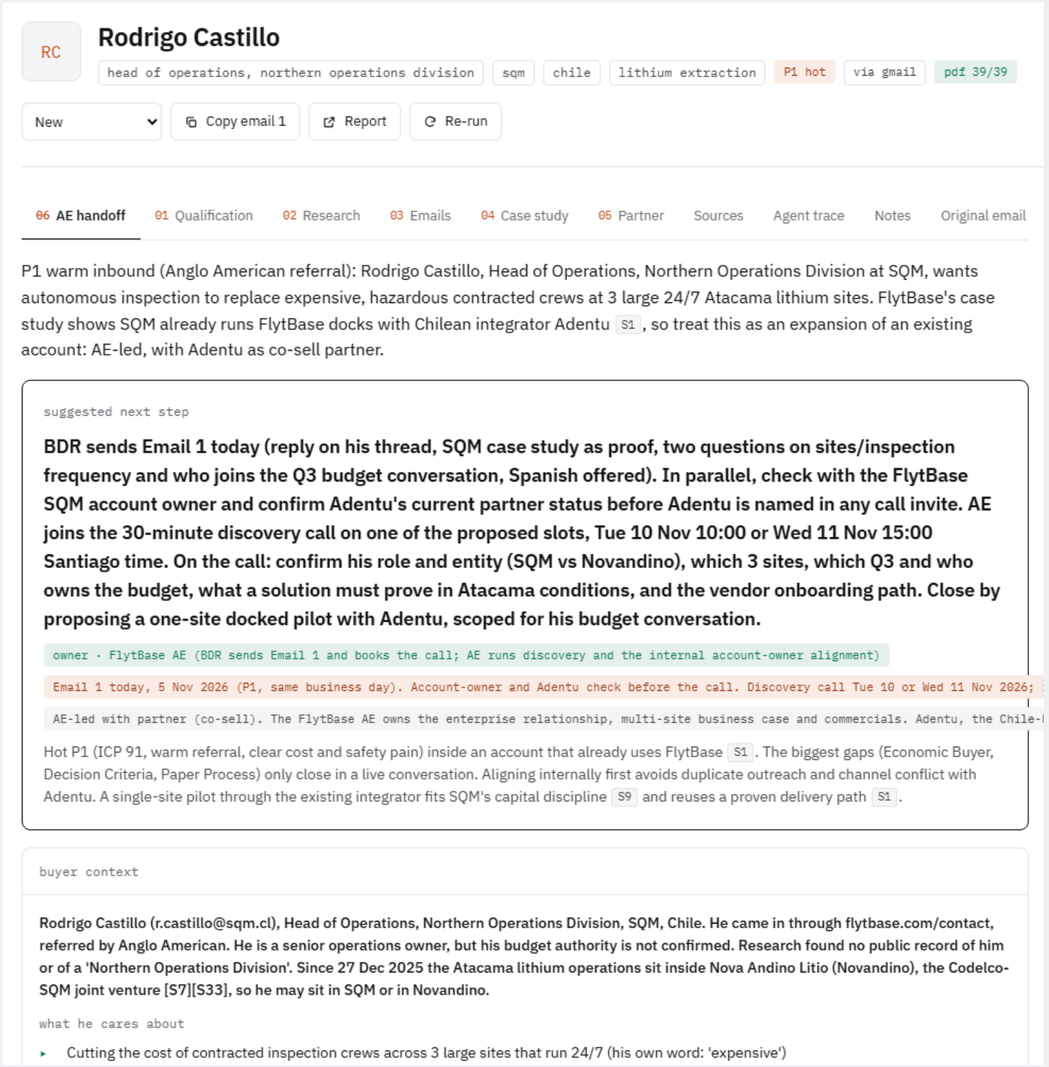
stage 06 · ae handoff

The AE knows exactly what to do next

Buyer context, qualification status, top 3 cited findings, case study and a next step with an owner and a date.

next step

Email 1 today. 30-min discovery call Tue 10 Nov 10:00 or Wed 11 Nov 15:00 Santiago. Confirm the account owner and Adentu first. Close on a one-site pilot.



live 2 · inspect the run

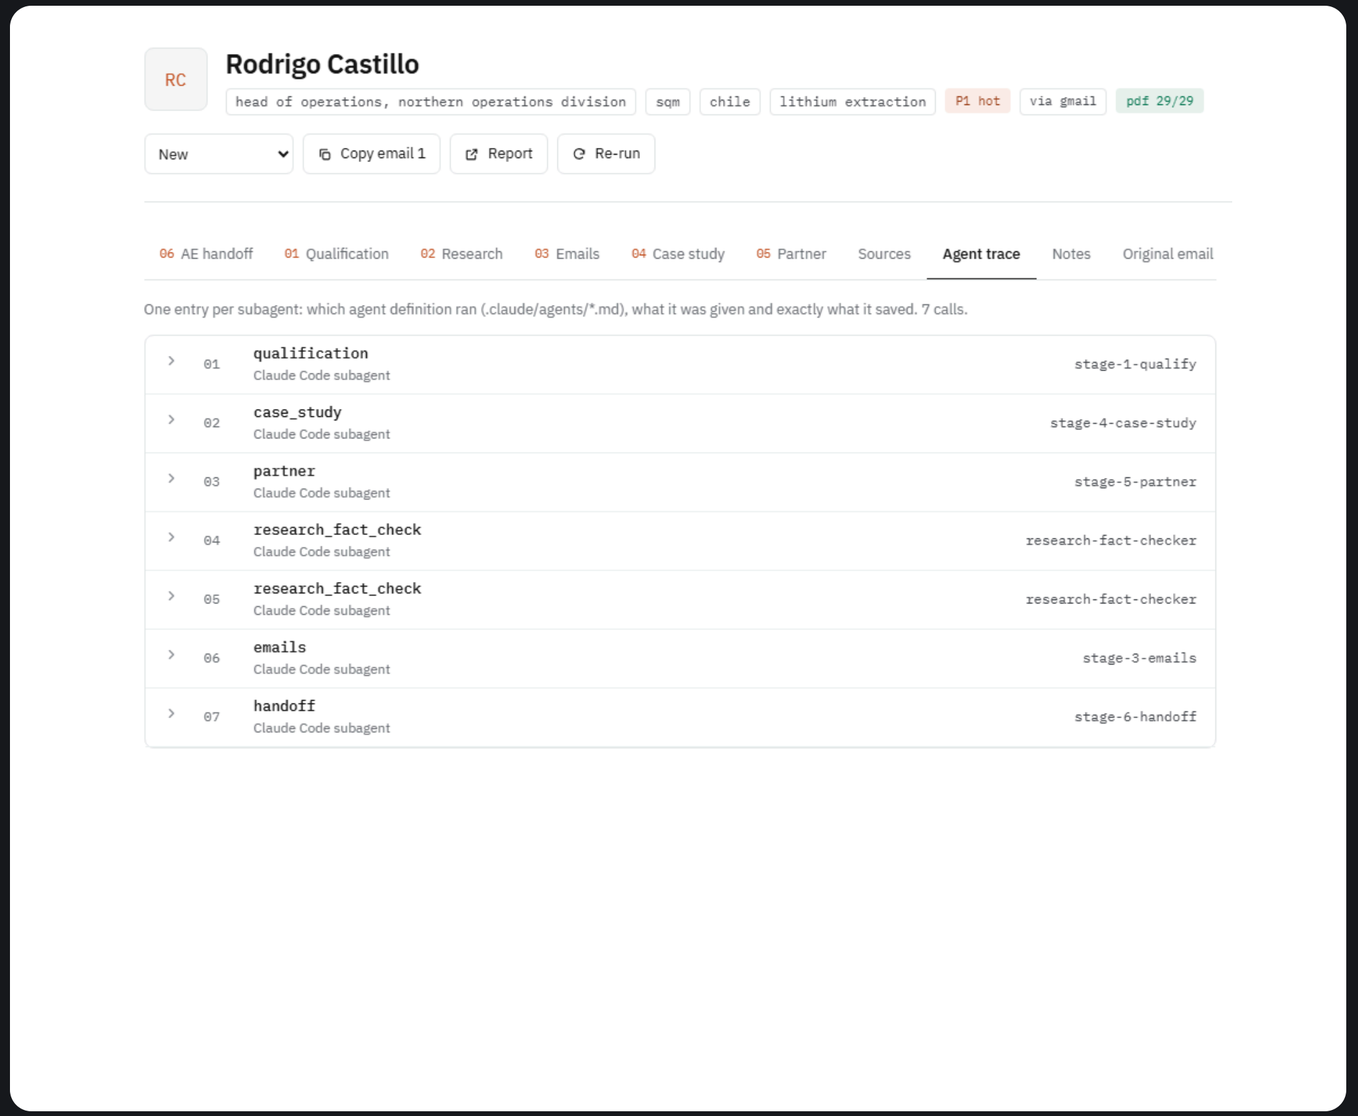
Every claim, check and call is inspectable

39/39
PDF + playbook checks

33
sources, raw text

Each stage tab ends with its PDF checklist. The Sources tab links every S#. The Agent trace shows which agent file ran, on what, and what it saved. REPORT.html, REPORT.pdf and one JSON per stage are in the repo.

show: checklist panel · sources · trace · report



it generalises

A second lead, a different answer

	SQM · Chile (the assignment)	E.ON · Germany (test lead)
Tier	P1 hot · 77 (Anglo American referral)	P2 warm · 67 (no referral, fit 74)
Research	Novandino JV is the likely buyer	Bayernwerk Netz is the likely buyer
Case study	SQM’s own deployment with Adentu	EnBW (solar, said plainly: not substations)
Partner	Adentu co-sell, Dronestore backup	Droniq co-sell, BVLOS approval
Emails	Santiago times, Spanish offer	CET times, German offer
Checks	39 / 39	29 / 29

Same code, no lead-specific logic. The score honestly put E.ON one point under P1, and the fact-checker removed 6 claims it couldn’t read from the source.

the design decision i’m proud of

The AI reads and writes. Python decides.

Scores

AI rates with evidence; Python computes the priority and tier.

Facts

Only stored sources can be cited; a separate agent checks; Python removes.

Names

A partner is only recommended if its name is in a source.

The brief

Every PDF bullet is a check that runs before a stage is saved.

---

16 → 0

The checker first saw only each page’s opening and removed 16 true claims. I found it, fixed the tool, re-checked: 0 removed.

6 cut

On E.ON, a Word file was stored unreadable. The checker cut every claim resting on it instead of trusting it.

limitations, honestly

What I’d improve with more time

Limitation	Why it matters	Next fix
The scheduler lives in the session	The cron stops when Claude Code closes and expires after 7 days	Headless runs on a real scheduler, or a Gmail push webhook
Judgement is still model output	Verdicts, ratings and wording can be off even when bounded	Slack approval before email 1; human review stays
Public web only	The contact and his division weren’t in any public source	CRM and enrichment data; LinkedIn through the browser
No CRM sync yet	Results live in MongoDB and the Lead Desk	Write leads and tasks to the team’s CRM

More than an email generator: a system that does the BDR’s first hour, every time.

System access: github.com/samruddhi-sontakke/flytbase-inbound-lead-agent

Live run output: outputs/run_2026-11-05_19-32-11 (REPORT.md · REPORT.pdf · one JSON per stage)

This deck: docs/Inbound-Lead-Desk-Walkthrough.pdf

Thank you · Samruddhi Sontakke

samruddhisontakke.com · sontakkesamruddhijain89@gmail.com · +91 72196 00321 · LinkedIn